

Seminar Report: Chatty

Pedro Garcia Mota, Pau Espin Pedrol

March 6, 2013

1 Introduction

Introduce in a couple of sentences the seminar and the main topic related to distributed systems it covers.

Implementation of a distributed system that will allow the user to chat among buddies. The main topic is communication between processes in different machines.

2 Work done

Present the new code that you have added or how you have implemented a required functionality by using small Erlang code snippets (you do not need to copy&paste all the code).

All the changes are commented in the code files.

3 Experiments

Describe the experiments you did and provide evidence of the results you got (e.g., use screenshots). In addition, you may provide figures or tables with experimental results of the system evaluation. For each seminar, we will provide you with some guidance on which kind of evaluation you should do.

On the first image we can see the server running on the top terminal window and two clients running on the bottom tabs.

4 Open questions

Try to answer all the open questions in the documentation. When possible, do experiments to support your answers.

Open Questions - 1.

- i) Does this solution scale when the number of users increase?

No, it would result in massive workload on the single server, and if the server crashes the whole system stops working.

- ii) What happens if the server fails?

All the service goes down. All users are unable to communicate each other.

- iii) Are the messages from a single client to any other client guaranteed to be delivered in the order they were issued? (hint: search for the 'order of message reception' in Erlang FAQ)

Yes, but only within one process, which is the case. According to Erlang's FAQ: "If there is a live process and you send it message A and then message B, it's guaranteed that if message B arrived, message A arrived before it."

- iv) What about the messages from different clients (hint: think on several clients sending messages concurrently)?

It's not guaranteed, as each client spends an extra time sending its information through the server, and there's the possibility that one client, because of network problems or high workload on the node for example, takes more time to send a message to the server than another client which actually issued it after the first one.

- v) What happens if a user joins/leaves the chat while the server is broadcasting a message?

It is undefined, it depends on the order of the events. It depends on the client ID being added to the list of clients before or after the server accesses it to send a message to all of the clients in the list.

Open Questions - 2.

- i) What happens if a server fails?

The clients connected to that server would no longer be able to chat with

other users.

ii) Are the messages from a single client to any other client guaranteed to be delivered in the order they were issued?

The behaviour is the same from the single server implementation.

iii) What happens if there are concurrent requests from servers to join or leave the system?

Nothing, as the FAQ states: "Delivery is guaranteed if nothing breaks"

iv) What are the advantages and disadvantages of this implementation regarding the previous one?

The main advantage is that the system is much more robust, if a server crashes the chat still works for the clients connected in the other servers. It is also much more scalable. On the other hand, the main disadvantages are the increased complexity of the system and the need of further resources for running the system.

5 Personal Opinion

Your personal opinion of the seminar, including whether it should be included in next year's course or not.

The seminar presents a easy introduction into the erlang language and concepts as message passing between different processes in the same machine and on different machines over the same network.